


Real cluster (results_k8s_v2, cpu workload): LLM vs. rule-based autoscalers

REAL K8S CLUSTER — measured, not simulated

Robustness — % of time P90 < 200 ms



Cost — @ \$0.04/vCPU-hr, one service

