



Real cluster (results_k8s_v2, io workload): LLM vs. rule-based autoscalers

REAL K8S CLUSTER — measured, not simulated



LLM (frozen) HPA KEDA RL (PPO) RL (DQN)