

Real cluster (results_k8s_v2, cpu workload): P50 → P90 → P99 latency spread per config

